

E6156 – Topics in SW Engineering: Cloud Computing

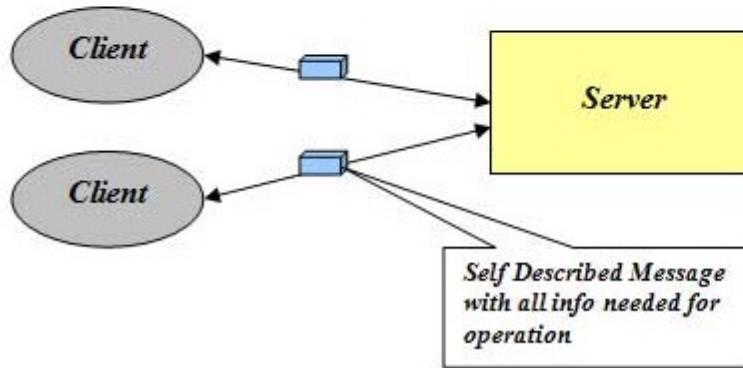
Lecture 8: Cool Stuff



Stateless, Sessions Example

Stateless

The notion of statelessness is defined from the perspective of the server. The constraint says that the server should not remember the state of the application. As a consequence, the client should send all information necessary for execution along with each request, because the server cannot reuse information from previous requests as it didn't memorize them. All info needed is in message.



Statelessness

2. Application State vs Resource State

It is important to understand the difference between the application state and the resource state. Both are completely different things.

Application state is server-side data that servers store to identify incoming client requests, their previous interaction details, and current context information.

Resource state is the current state of a resource on a server at any point in time – and it has nothing to do with the interaction between client and server. It is what we get as a response from the server as the API response. We refer to it as resource representation.

REST statelessness means being free from the application state.

- This concept can be confusing with experience with other approach to session state.
- Some application frameworks did/do keep session state on the server.
- Application servers and frameworks have various support for “session state,” e.g. <https://flask-session.readthedocs.io/en/latest/>

3. Advantages of Stateless APIs

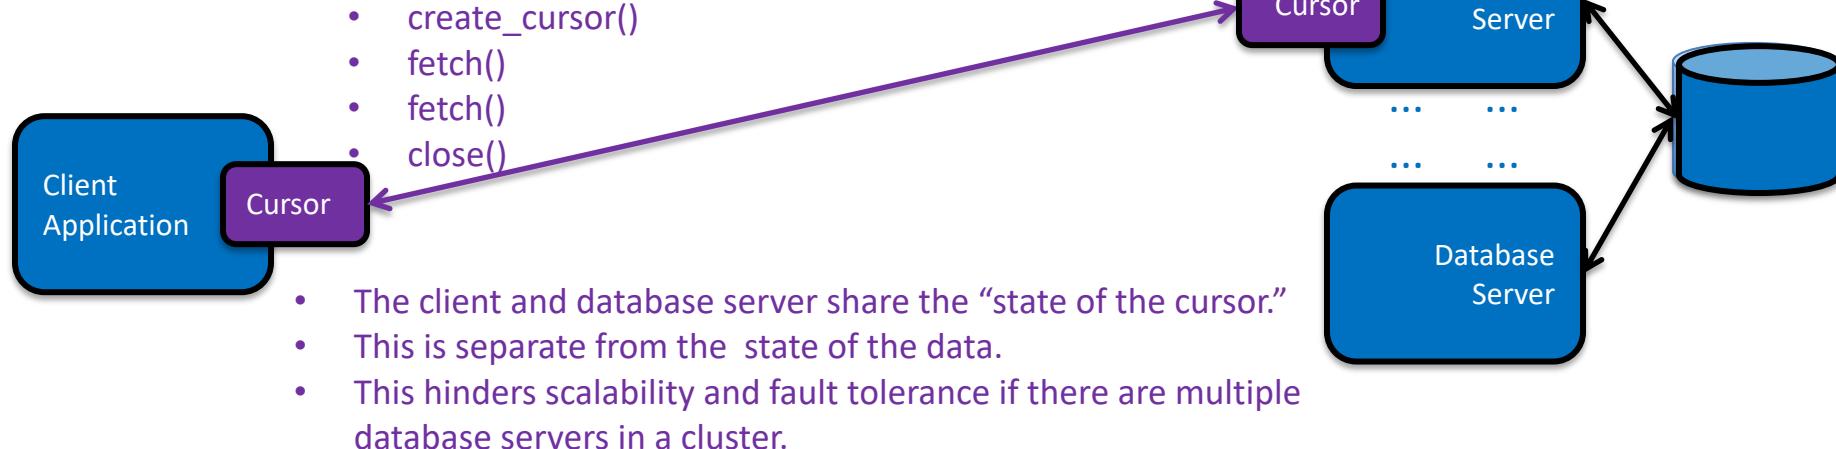
There are some very noticeable advantages of having REST APIs **stateless**.

1. Statelessness helps in **scaling the APIs to millions of concurrent users** by deploying it to multiple servers. Any server can handle any request because there is no session related dependency.
2. Being stateless makes REST APIs **less complex** – by removing all server-side state synchronization logic.
3. A stateless API is also **easy to cache** as well. Specific softwares can decide whether or not to cache the result of an HTTP request just by looking at that one request. There's no nagging uncertainty that state from a previous request might affect the cacheability of this one. It **improves the performance** of applications.
4. The server never loses track of “where” each client is in the application because the client sends all necessary information with each request.

Simple Example from Databases

- “In computer science, a database cursor is a mechanism that enables traversal over the records in a database. Cursors facilitate subsequent processing in conjunction with the traversal, such as retrieval, addition and removal of database records. The database cursor characteristic of traversal makes cursors akin to the programming language concept of iterator.”

([https://en.wikipedia.org/wiki/Cursor_\(databases\)](https://en.wikipedia.org/wiki/Cursor_(databases)))



Some Code

```
def t_cursor():

    print("\nUsing cursors ...")
    sql = "select * from users";

    for i in range(0,5):
        if i == 0:
            res = with_cursor(done=False, sql=sql)
        elif i == 4:
            res = with_cursor(done=True)
        else:
            res = with_cursor()

    print("Res = ", res)
```

```
def t_cursor():

    print("\nUsing cursors ...")
    sql = "select * from users";

    for i in range(0,5):
        if i == 0:
            res = with_cursor(done=False, sql=sql)
        elif i == 4:
            res = with_cursor(done=True)
        else:
            res = with_cursor()

    print("Res = ", res)
```

```
def t_without_cursor():

    print("\nNot using cursors.")

    sql = "select * from users";

    for i in range(0,5):
        res = without_cursor(sql=sql, offset=i)
        print("Res = ", res)
```

```
def without_cursor(sql, offset):

    cursor = conn.cursor()
    tmp_sql = sql + " limit 1" + " offset " + str(offset)
    res = cursor.execute(tmp_sql)
    res = cursor.fetchone()
    cursor.close()
    return res
```



Calling APIs

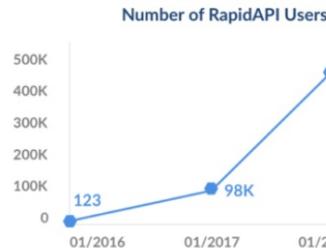
A World of APIs



Enterprise API Hub Add Your API

Blog » APIs » Top 50 Most Popular APIs on RapidAPI (2022)

RAPIDAPI GROWTH



APIs in the RapidAPI Marketplace

Over
8,000
Up from 200 APIs as of 11/2016



README.md

Public APIs

A collective list of free APIs for use in software and web development

Status

Number of Categories 51 Number of APIs 1425

Tests of push & pull failing Validate links failing Tests of validate package passing

The Project

[Contributing Guide](#) • [API for this project](#) • [Issues](#) • [Pull Requests](#) • [License](#)



Explore the World of APIs

Browse the largest network of APIs, workspaces, and collections by developers across the planet

Explore

Categories

In the spotlight

Browse

Teams

1.8M

Workspaces

100.1k

APIs

13.4k

Collections

208.0k

Popular on Postman

Discover what's new and popular on the Postman API Network!

[View all →](#)

Stripe API [09-29-2022]

StripeDev

This is a reference collection for the Stripe API from 09-29-2022. Fork it to your workspace to try out various Stripe API requests.

Fork | 400+ Watch | 222

Meraki Dashboard API - v1.13.0

Cisco Meraki

The primary API for interacting with a Cisco Meraki network.

Fork | 1k+ Watch | 472

Braze Endpoints

Braze

April 22, 2022 Datadog API Collection

Datadog

This is the collection for the Datadog REST API, created on

Two APIs I Use in My Project

smarty

Products Pricing Resources Contact |  Cart Log Out Account →

← Demos

US Street Address API

US Street Address API

US ZIP Code API

US Autocomplete Pro API

US Extract API

International Street Address API

International Address Autocomplete API

US Rooftop Geocoding

US Reverse Geocoding

Canvas LMS - REST API and Extensions Documentation

Basics

GraphQL

API Change Log

SIS IDs

Pagination

Throttling

Compound Documents

File Uploads

API Endpoint Attributes

Masquerading

OAuth2

OAuth2 Overview

OAuth2 Endpoints

Developer Keys

Resources

All Resources Reference

Welcome to the Canvas LMS API Documentation

Canvas LMS includes a REST API for accessing and modifying data externally from the main application, in your own programs and scripts. This documentation describes the resources that make up the API.

To get started, you'll want to review the general basics, including the information below and the page on [Authentication using OAuth2](#).

API Changes

For API resources, such as the API Change Log for additions, changes, deprecations, and removals, view the [Canvas API page](#) in the Canvas Community.

API Terms of Service

Please carefully review [The Canvas Cloud API Terms of Service](#) before using the API.

Schema

All API access is over HTTPS, against your normal Canvas domain.

All API responses are in [JSON format](#).

All integer IDs in Canvas are 64-bit integers. String IDs are also used in Canvas.

To force all IDs to strings add the request header `Accept: application/json+canvas-string-ids`. This will cause Canvas to return even integer IDs as strings, preventing problems with languages (particularly JavaScript) that can't properly process large integers.

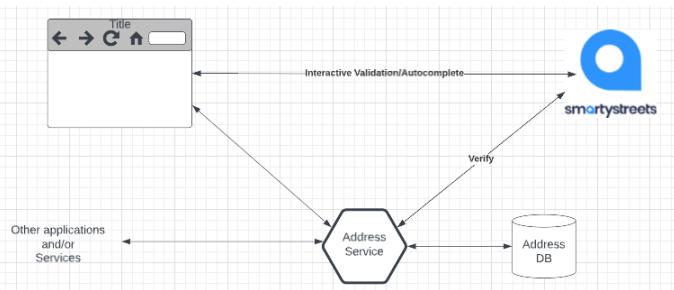
- Addresses:

- People are really bad at entering mailing addresses.
 - 520 w 120th st NYC
 - 520 West 120 St New York, NY
 -
- We want a canonical, verified representation.

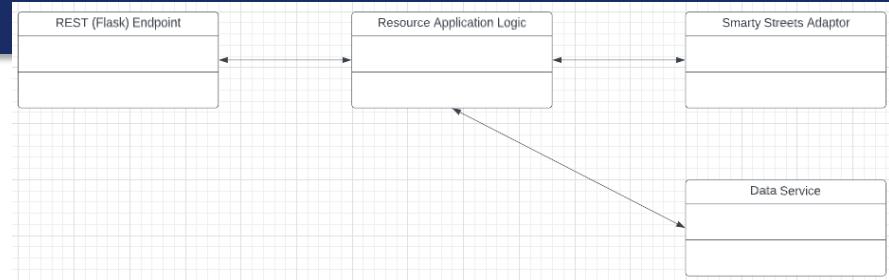
- CourseWorks

- Students will access my web application to set up projects, teams, ...
- I want to verify that students are in the class, and keep list of students up to date.

Address Example

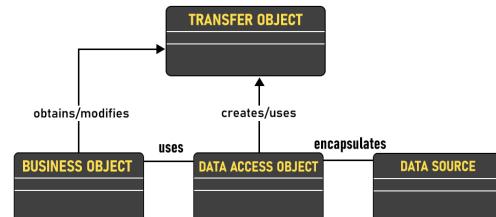


- Use autocomplete functions to help users correctly enter addresses
<https://jsfiddle.net/smartystreets/j1Lq5say>
- But, ...
 - A microservice cannot assume that a valid UI is calling the service.
 - It must support any UI or endpoint authorized to call the service.
 - So, we will also validate in the microservice implementation.
 - Example code walkthrough and online tools.



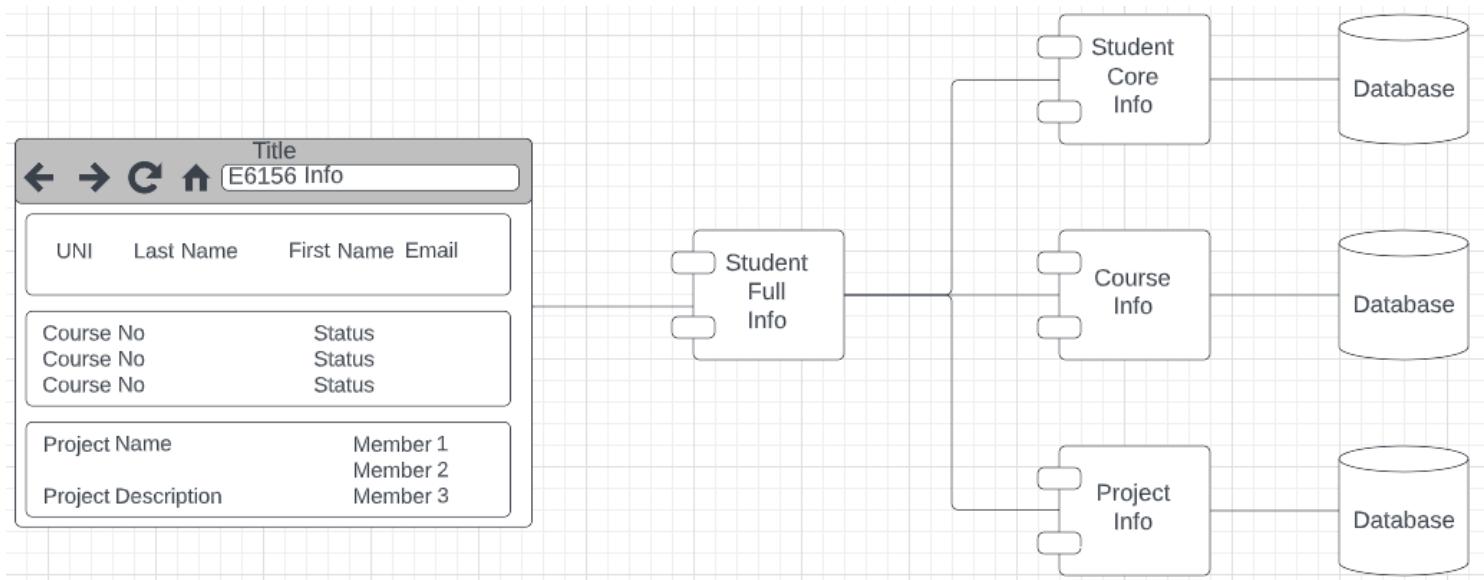
Microservice Implementation

1. Receive JSON body in PUT/POST.
2. Call API to verify.
3. If verified, save in DB.
4. URL path is the delivery point
 1. Zipcode
 2. Zip plus 4
 3. Delivery point



Synchronous versus Asynchronous

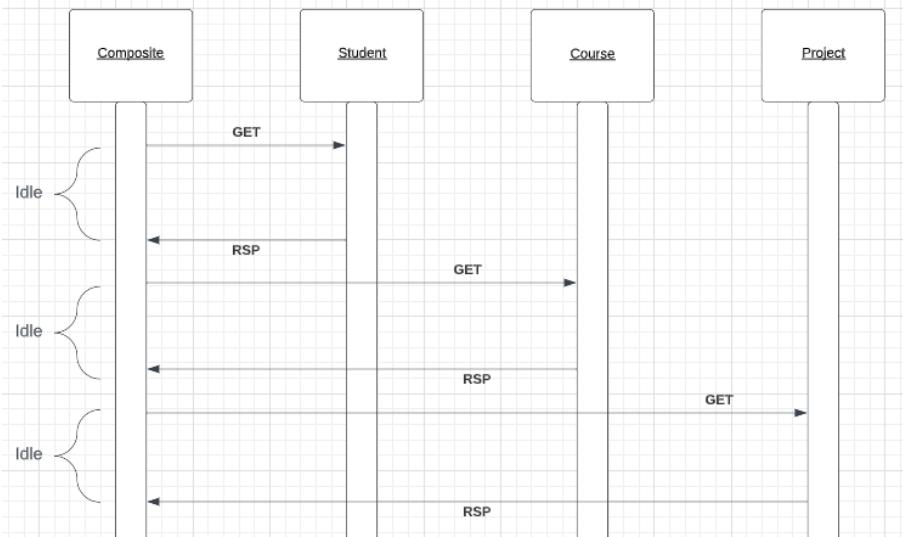
Synchronous versus Asynchronous



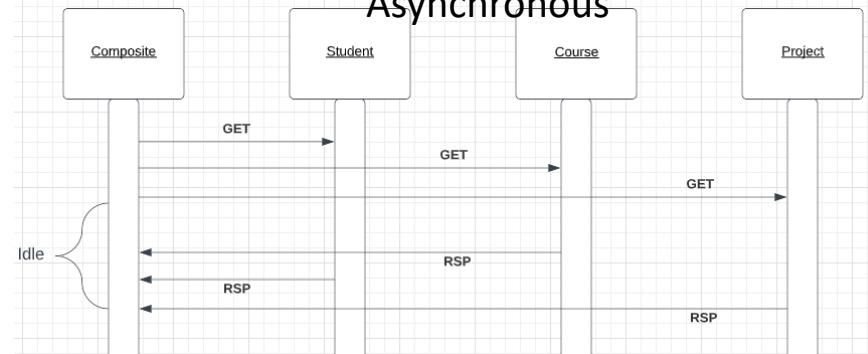
- CRUD on composite service `StudentFullInfo` requires CRUD on 3 microservices.
- There are two styles for calling the component (composed services)
 - Synchronous, one after another, waiting each time.
 - Asynchronous: Call all three in parallel, and "gather the answer."

Synchronous versus Asynchronous

Synchronous



Asynchronous



Show code from:

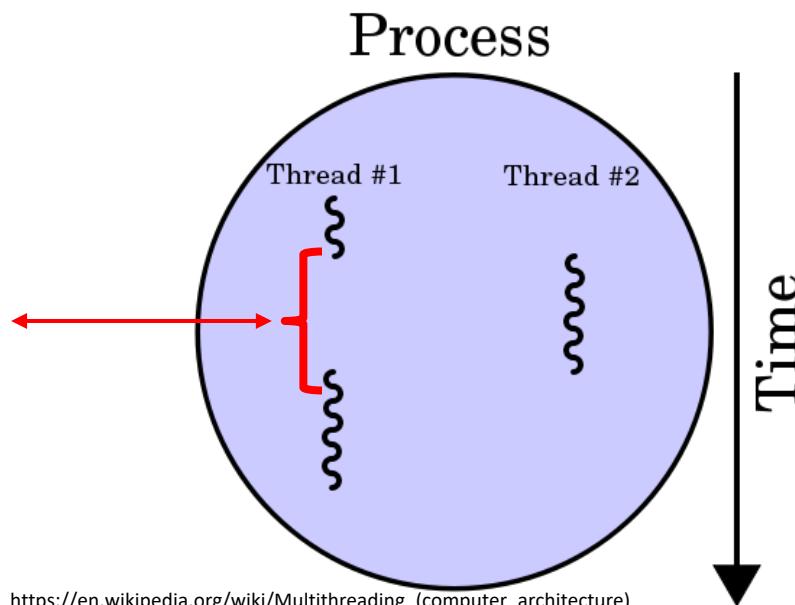
- `call_urls_synch_asynch`
 - `async3`
 - `async_p_mysql`
- Async-python
- `av_requests`
 - `av_async_run`

- The elapsed time for calling synchronous is *the sum of the calls' response times*.
- The elapsed time for calling asynchronous is the maximum response time of the 3 calls.
- Asynchronous:
 - Has better elapsed times and less “application is idle waiting” time, but
 - The logic is more complex, especially if the composite services has to “undo” some calls because others failed.

Concurrency

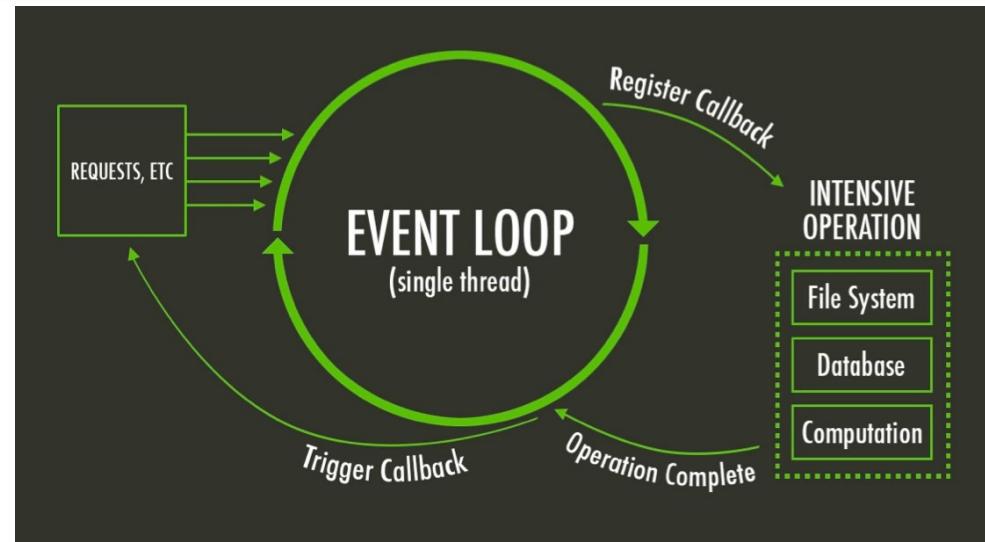
- A running program executes in an operating system *process*. There are many approaches to concurrency:
 - Multiprocessing:
 - The system has several “active processes.”
 - Only one process is executing on a CPU (core) as a time.
 - The OS switches between processes when the running process performs a blocking operations, e.g. file IO, remote network call,
 - Multithreading:
 - There are several “active threads” within a process.
 - (Usually) Only one thread is executing on a CPU (core) at a time.
 - When the thread performs a blocking operation, the process continues to run but executes a different, ready thread.
 - Coroutines:
 - There is a single active thread within a process.
 - When a thread performs a blocking operation, the programming language runtime does weird internal stuff
 - Makes a copy of the current coroutines environment (call stack,)
 - Starts a ready coroutine by switching to its call stack and next statement.
 - The thread and process do not block.
- A *context switch* is changing the currently running process, thread or task.
 - There is a performance overhead for a context switch.
 - Overhead(multiprocessing) >> Overhead(multithreading) >> Overhead(coroutines)

Threads, Promises, Asynch Functions,



[https://en.wikipedia.org/wiki/Multithreading_\(computer_architecture\)](https://en.wikipedia.org/wiki/Multithreading_(computer_architecture))

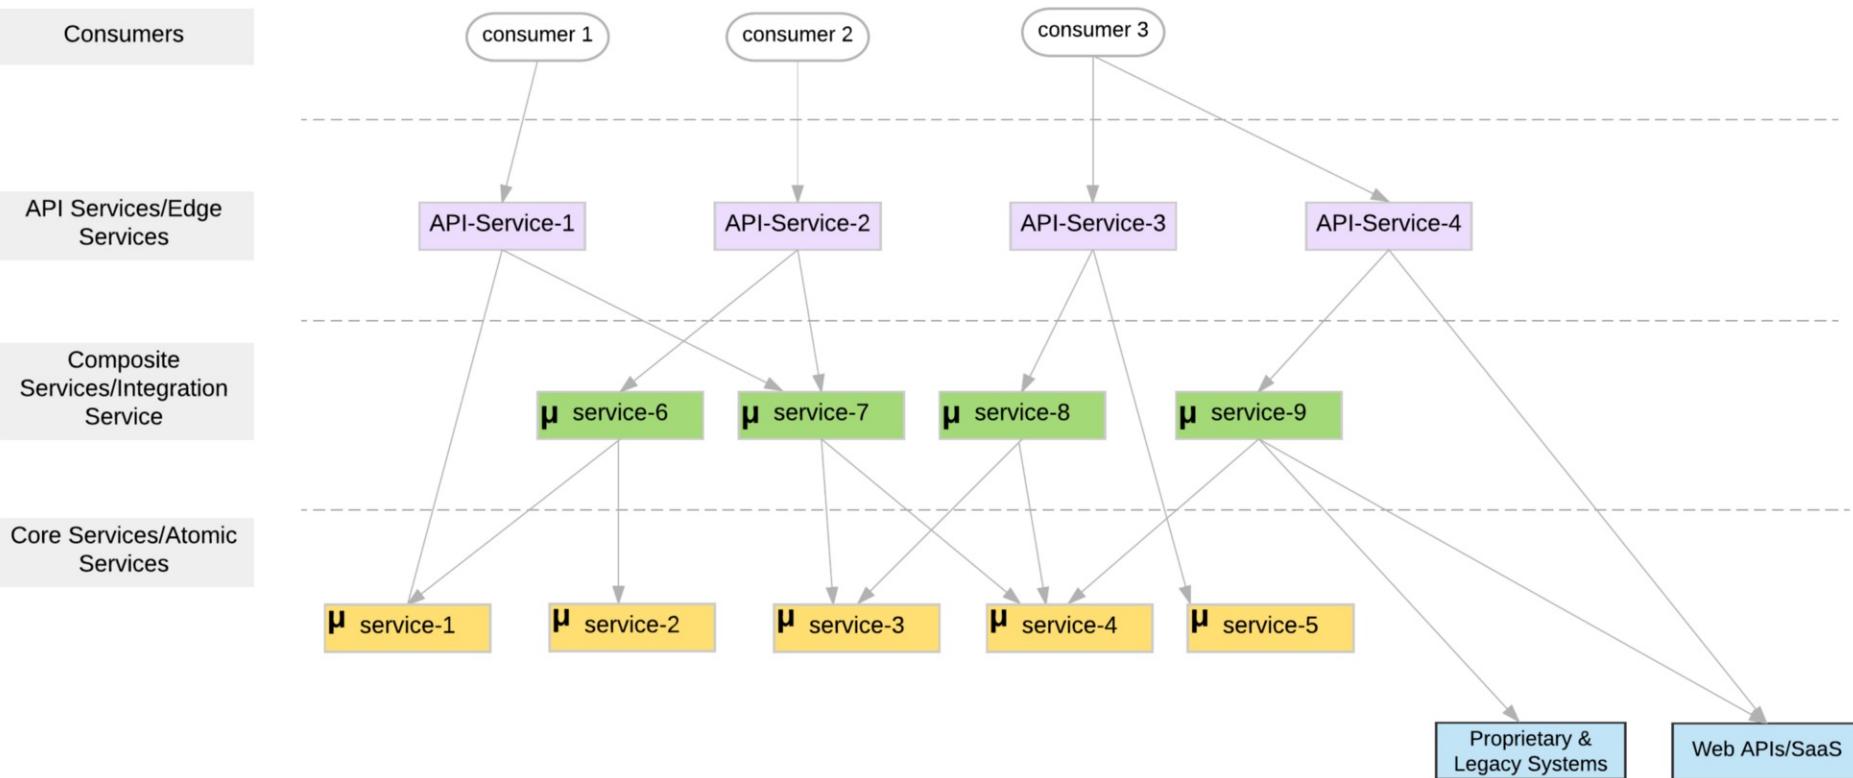
- Multithreading
 - Thread #1 “blocks” for an IO or API call.
 - Process switches to “ready” thread #2



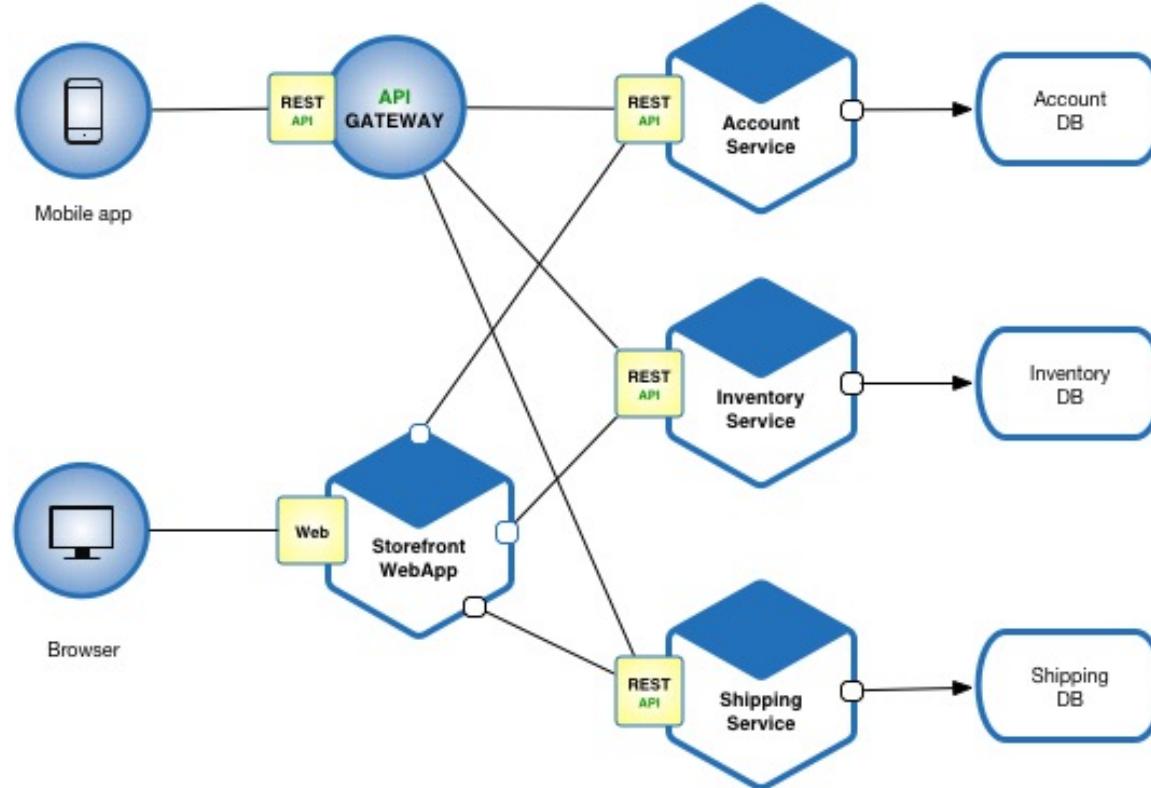
- **Callbacks**
 - Incoming requests map to an event.
 - There is a single thread running application code associated with an event..
 - Long running calls get assigned to a separate worker thread, and register a callback.
 - Completion of an operation results in an event, which resumes the application via a callback.

Service Composition

Microservice Layers

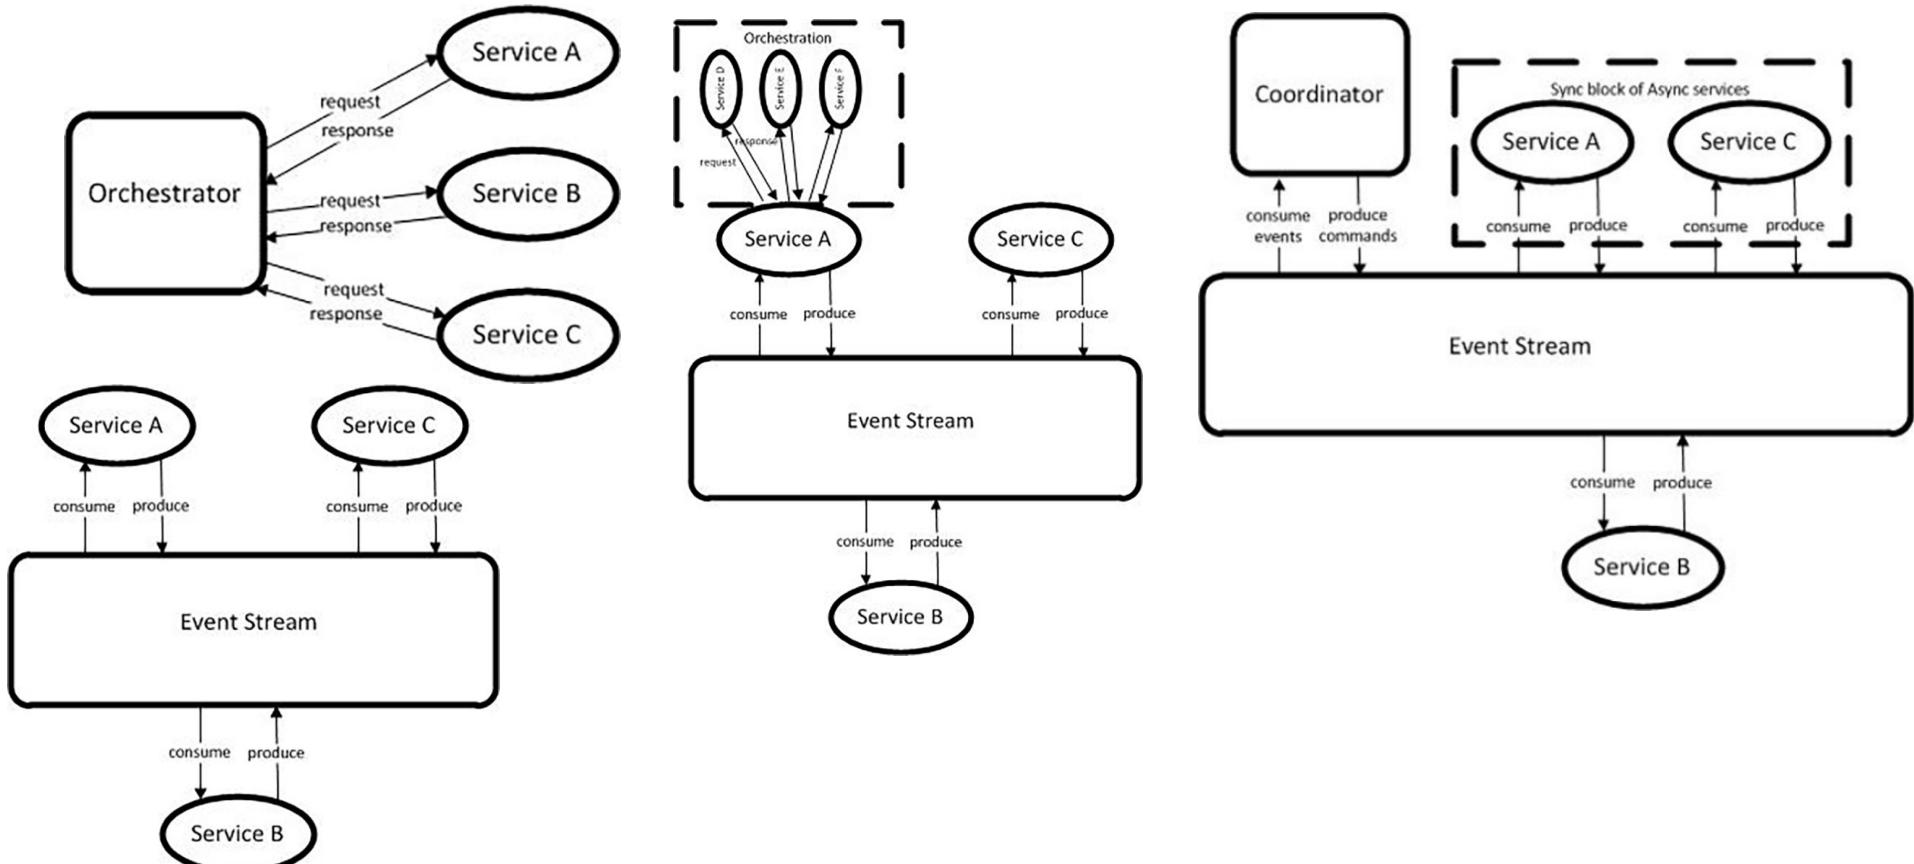


Another View



Models

(<https://medium.com/capital-one-tech/microservices-when-to-react-vs-orchestrate-c6b18308a14c>)

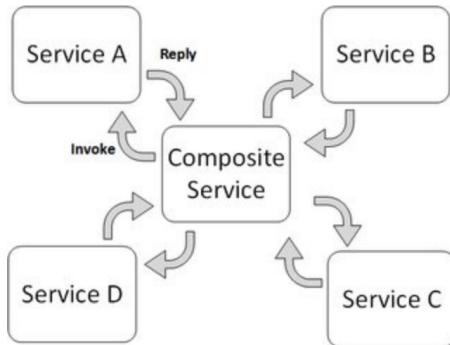


Concepts

Service orchestration

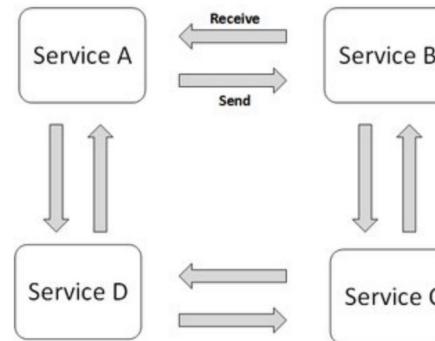
Service orchestration represents a single centralized executable business process (the orchestrator) that coordinates the interaction among different services. The orchestrator is responsible for invoking and combining the services.

The relationship between all the participating services are described by a single endpoint (i.e. the composite service). The orchestration includes the management of transactions between individual services. Orchestration employs a centralized approach for service composition.



Service Choreography

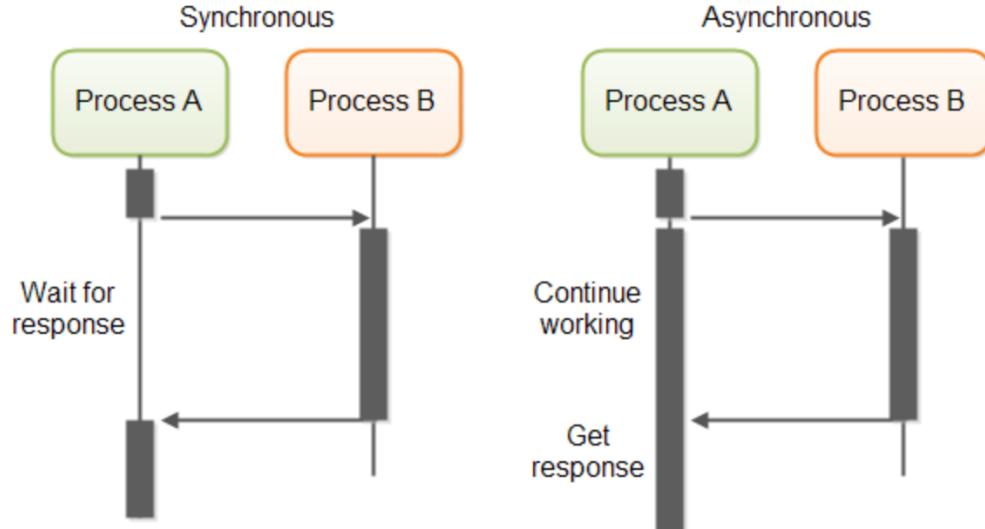
Service choreography is a global description of the participating services, which is defined by exchange of messages, rules of interaction and agreements between two or more endpoints. Choreography employs a decentralized approach for service composition.



The choreography describes the interactions between multiple services, whereas orchestration represents control from one party's perspective. This means that a **choreography** differs from an **orchestration** with respect to where the logic that controls the interactions between the services involved should reside.

These are not formally, rigorously defined terms.

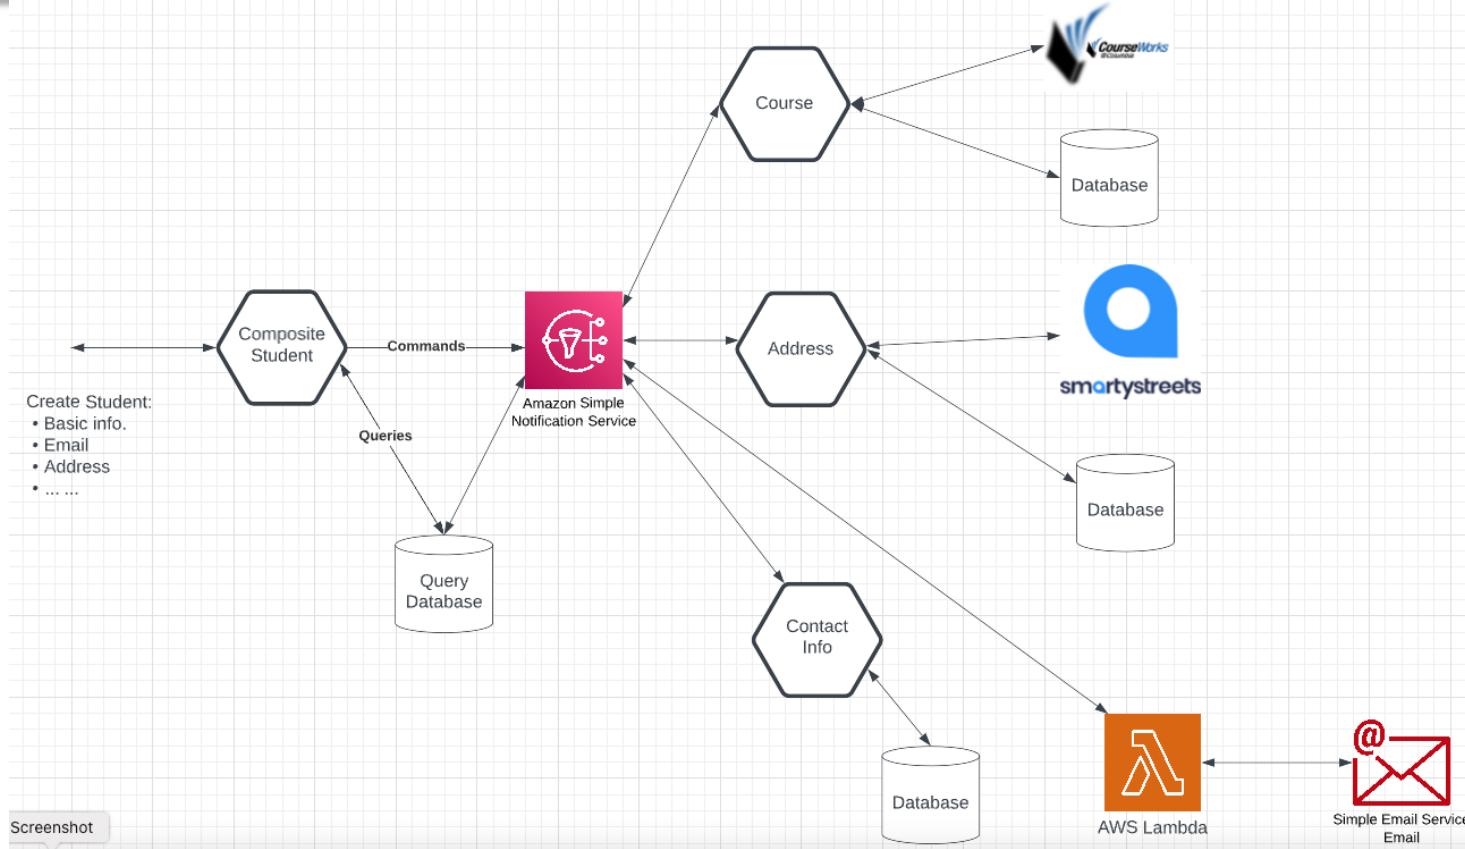
Service Orchestration/Composition



RESTful HTTP calls can be implemented in both a synchronous and asynchronous fashion at an IO level.

- One call to a composite service
 - May drive calls to multiple other services.
 - Which in turn, may drive multiple calls to other services.
- Synchronous (Wait) and calling one at a time, in order is
 - Inefficient
 - Fragile
 -
- Asynchronous has benefits but can result in complex code.

What I (We) Will Implement



Screenshot

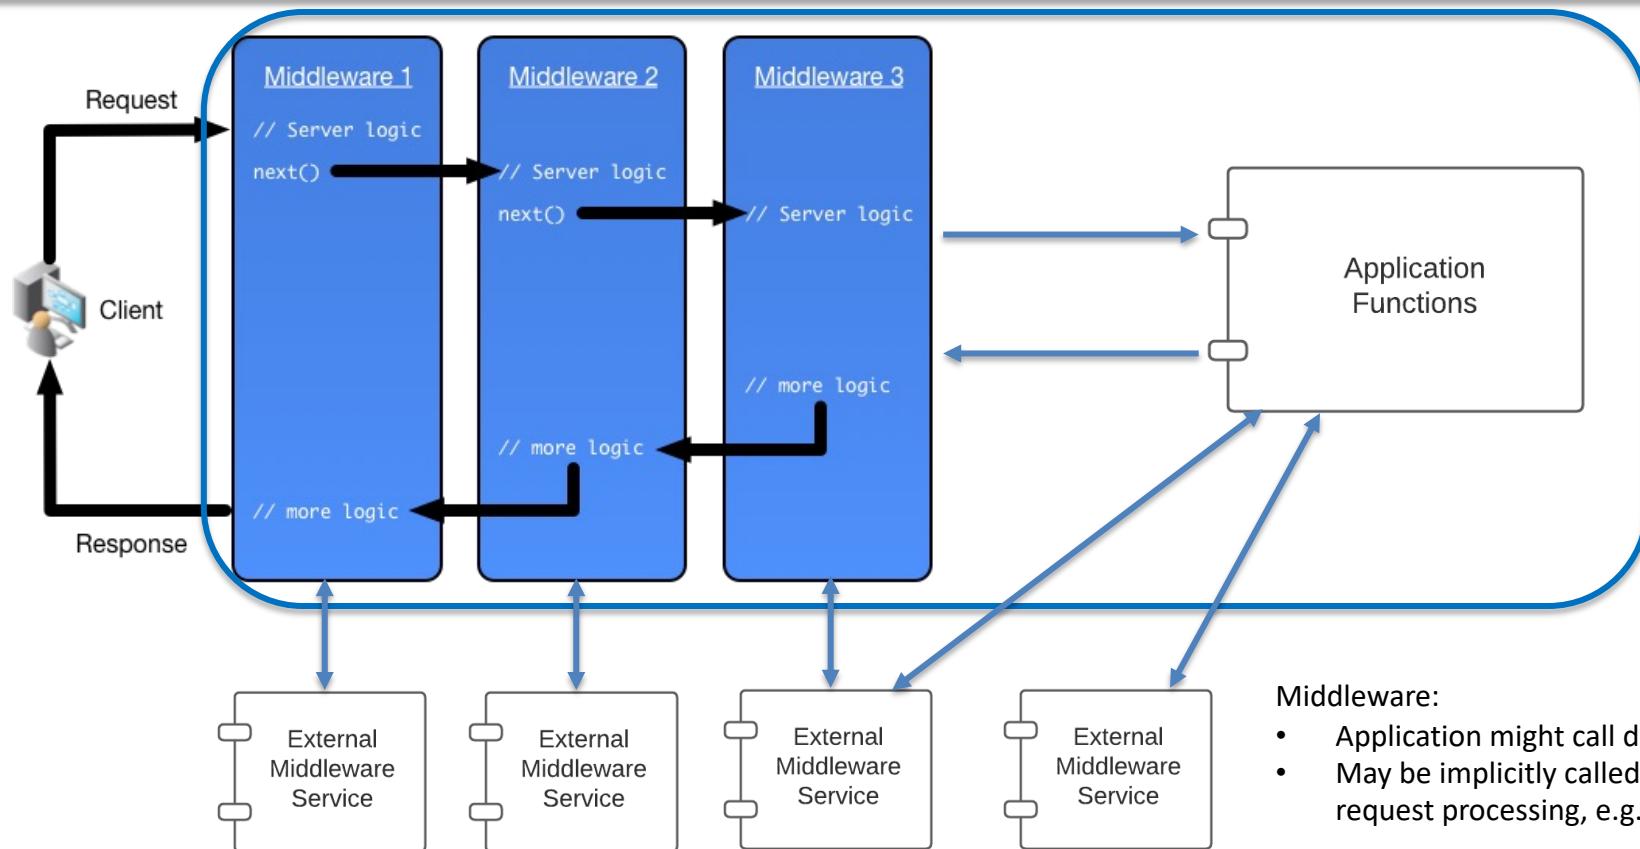
Middleware

What is middleware?

Middleware is software that lies between an operating system and the applications running on it. Essentially functioning as hidden translation layer, middleware enables communication and data management for distributed applications. It's sometimes called plumbing, as it connects two applications together so data and databases can be easily passed between the "pipe." Using middleware allows users to perform such requests as submitting forms on a web browser, or allowing the web server to return dynamic web pages based on a user's profile.

Common middleware examples include database middleware, application server middleware, message-oriented middleware, web middleware, and transaction-processing monitors. Each program typically provides messaging services so that different applications can communicate using messaging frameworks like simple object access protocol (SOAP), web services, representational state transfer (REST), and JavaScript object notation (JSON). While all middleware performs communication functions, the type a company chooses to use will depend on what service is being used and what type of information needs to be communicated. This can include security authentication, transaction management, message queues, applications servers, web servers, and directories. Middleware can also be used for distributed processing with actions occurring in real time rather than sending data back and forth.

Middleware – Conceptual Model



Middleware and Flask

```
# These methods get called before/after. You can check security information for all requests.
```

```
@application.before_request
```

```
def before_decorator():
```

```
    a_ok = sec.check_authentication(request)
```

```
    print("a_ok = ", a_ok)
```

```
    if a_ok[0] != 200:
```

```
        handle_error(a_ok[0], a_ok[1], a_ok[2])
```

- White list of URLs
- If not on white list:
 - Verify token
 - Verify authorization

```
@application.after_request
```

```
def after_decorator(rsp):
```

```
    print("... In after decorator ...")
```

```
    notify.notify(request, rsp)
```

```
    return rsp
```

- List of filters of the form:
 - URL
 - HTTP Method
 - SNS topic to use
- Send an event to SNS if matches filter

```
@application.route("/api/users", methods=["GET", "POST"])
```

```
def users_get():...
```

Show some demo code.

So, How Will We Use This?

- Add before_request and after_request to microservices.
(Note: There are other approaches).
 - Implement two middleware interceptors:
 - Security
 - Configured with a JSON file.
 - File lists which {path, method} not being logged on.
 - While IDs can perform which methods on which paths.
 - Automates driving login redirect (401 – Unauthorized)
 - Checks authorization (403– Forbidden)
 - Notification:
 - Configured with JSON in environment variables.
 - File lists which operations emit a before and/or after event to an SNS topic.
 - We will use notification to drive some other things like webhooks.
- 
- I will demo middleware and webhook today.
 - Your project will use middleware and SNS.

Cloud Identity (Open ID, OAuth2,) Reminder

/Users/donaldferguson/Dropbox/00NewProjects/google_login

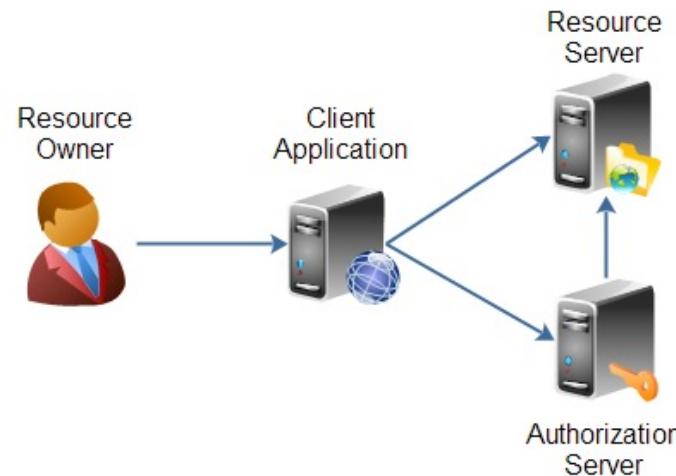
Overview (<http://tutorials.jenkov.com/oauth2/index.html>)

- Resource Owner
 - Controls *access* to the “data.”
 - Facebook **user**, LinkedIn **user**, ...
- Resource Server
 - The website that holds/manages info.
 - Facebook, LinkedIn, ...
 - And provides access API.
- Client Application
 - “The product you implemented.”
 - Wants to read/update
 - “On your behalf”
 - The data the Resource Server maintains
- Authorization Server
 - Grants/rejects authorization
 - Based on Resource Owner decisions.
 - Usually (logically) the same as Resource Server.

OAuth 2.0 defines the following roles of users and applications:

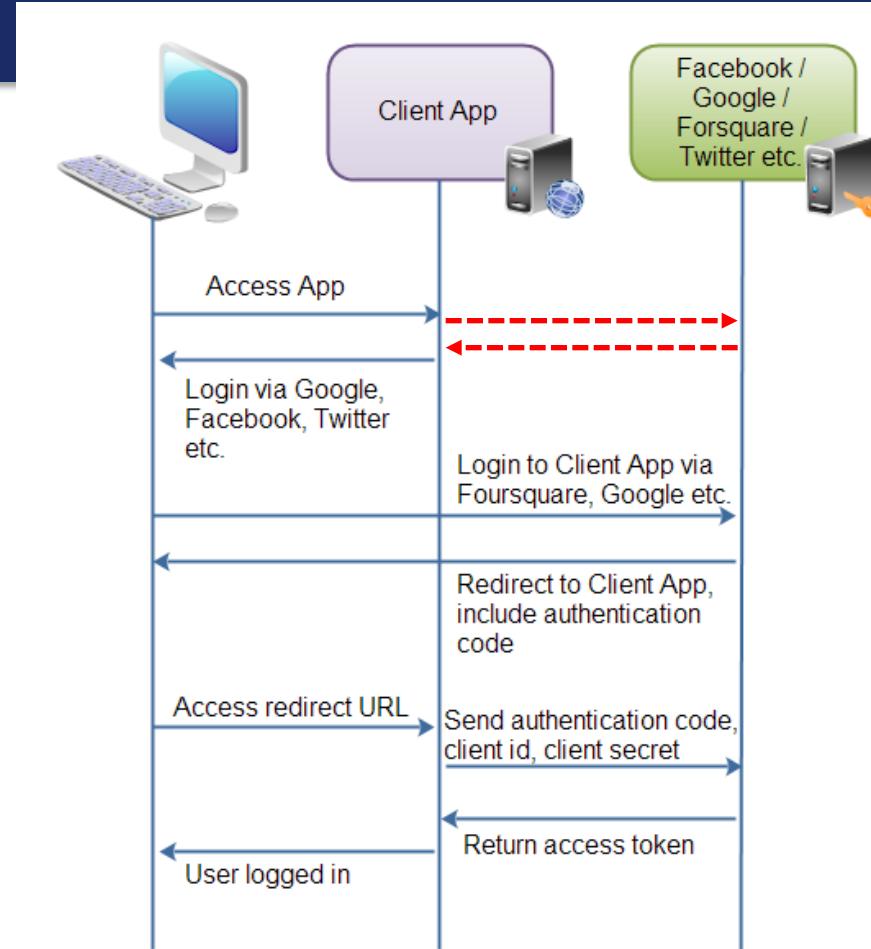
- Resource Owner
- Resource Server
- Client Application
- Authorization Server

These roles are illustrated in this diagram:

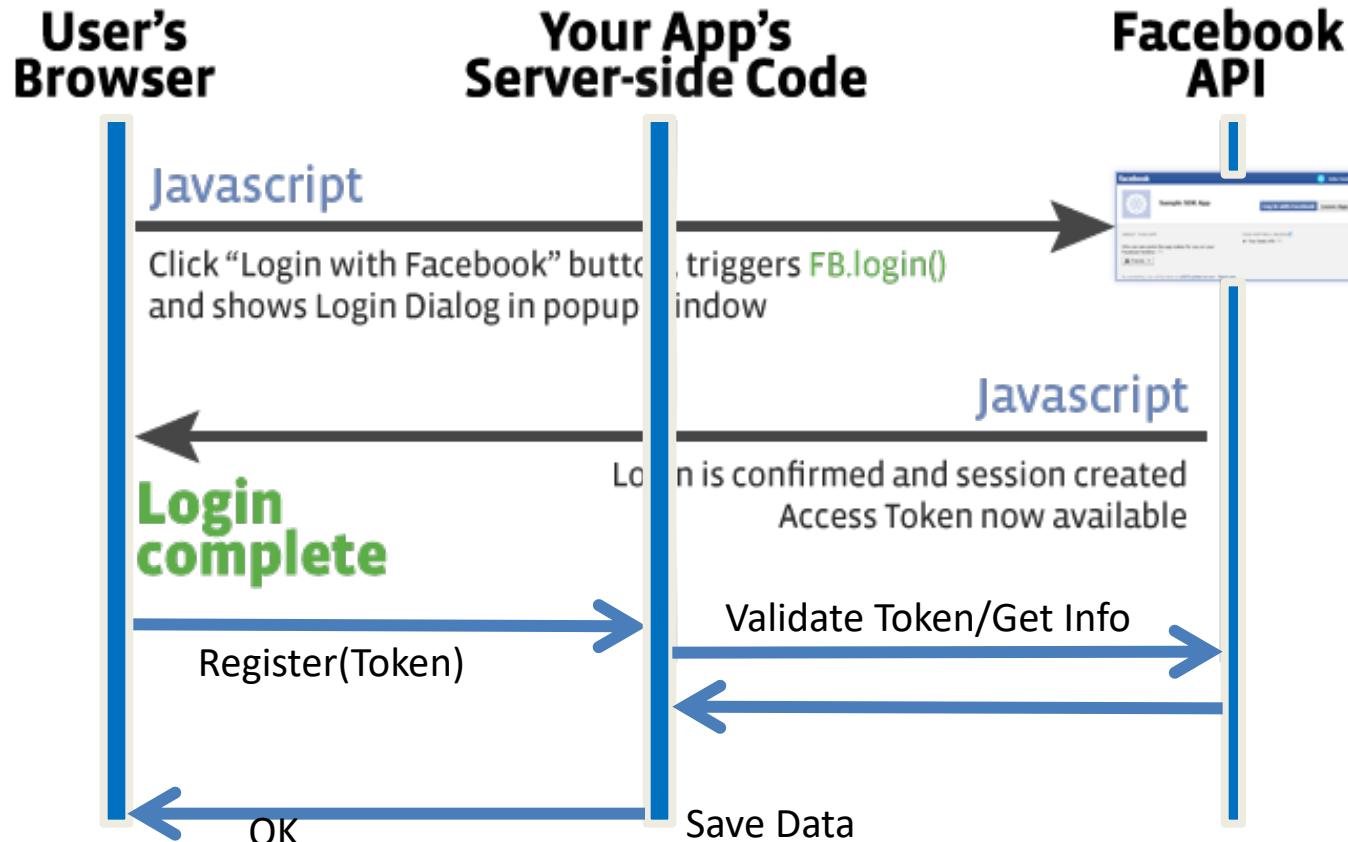


Roles/Flows

- User Clicks “Logon with XXX”
 - Redirect user to XXX
 - With Client App application ID.
 - And permissions app requests.
 - **Code MAY have to call Resource Server to get a token to include in redirect URL.**
- Browser redirected to XXX
 - Logon on to XXX prompt.
 - Followed by “do you grant permission to ...”
 - Clicking button drives a redirect back to Client App. URL contains a temporary token.
- User/Browser
 - Redirected to Client App URL with token.
 - Client App calls XXX API
 - Obtains access token.
 - Returns to User.
- Client App can use access token on API calls.

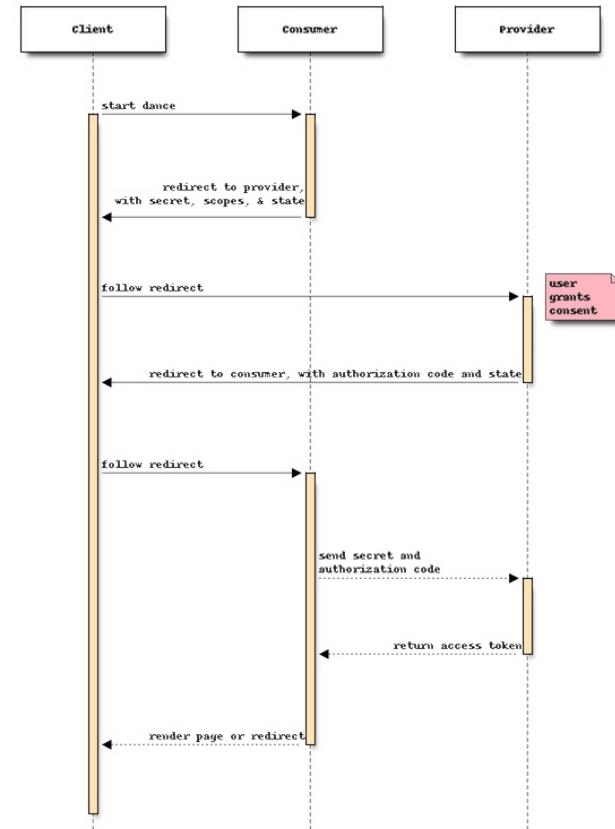


Facebook



Walkthrough: Middleware and Security

- Explain why I have to go incognito.
- Show:
 - Project code
 - GCP application registration page.
 - Token in the header (Postman)
- I normally do these things explicitly in my code.
 - I used Facebook in the past.
 - I decided to use Google because I expect people using my team management application to have Lionmail/Google accounts.
 - For expediency, I used Flask-Dance (<https://flask-dance.readthedocs.io/en/latest/>)
 - I understand the magic in general, but ...
 - I have not yet “learned the flask-dance magic” and how it uses the protocol under the hood.
- The net is:
 - OAuth2 is about *authorization*, but ...
 - Getting the basic profile proves *authentication*
 - And, I can use to simplify the registration process.



To Show

- Code walkthrough.
- Cookie/header.
- Certificate.
-

Event Drive Processing

Publish/Subscribe

Pub/Sub – Event Driven Processing

"In [software architecture](#), **publish–subscribe** is a [messaging pattern](#) where senders of [messages](#) do not program the messages to be sent directly to specific receivers, called subscribers, but instead categorize publishers by topic. Subscribers express interest in one or more topics and only receive knowledge of which subscribers, if any, there may be. Similarly, subscribers express interest in one or more topics and only receive messages that are of interest, without knowledge of which publishers, if any, there are."

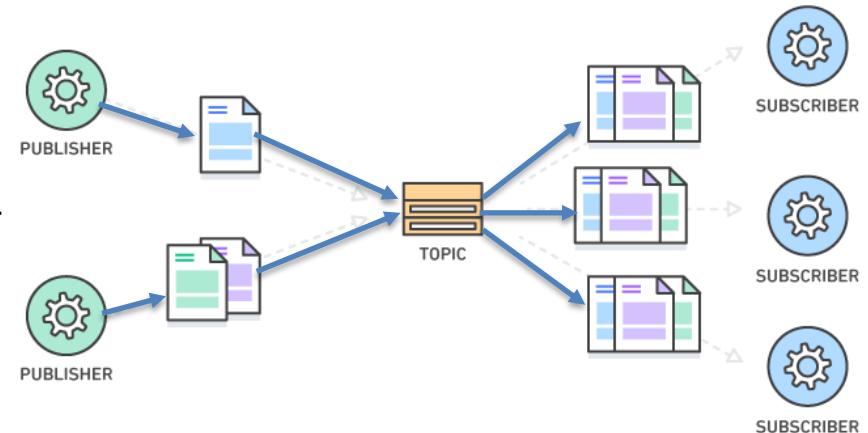
Publish–subscribe is a sibling of the [message queue](#) paradigm, and is typically one part of a larger [message-oriented middleware](#) system. Most messaging systems support both the pub/sub and message queue models in their [API](#), e.g. [Java Message Service](#) (JMS).

This pattern provides greater network [scalability](#) and a more dynamic [network topology](#), with a resulting decreased flexibility to modify the publisher and the structure of the published data."

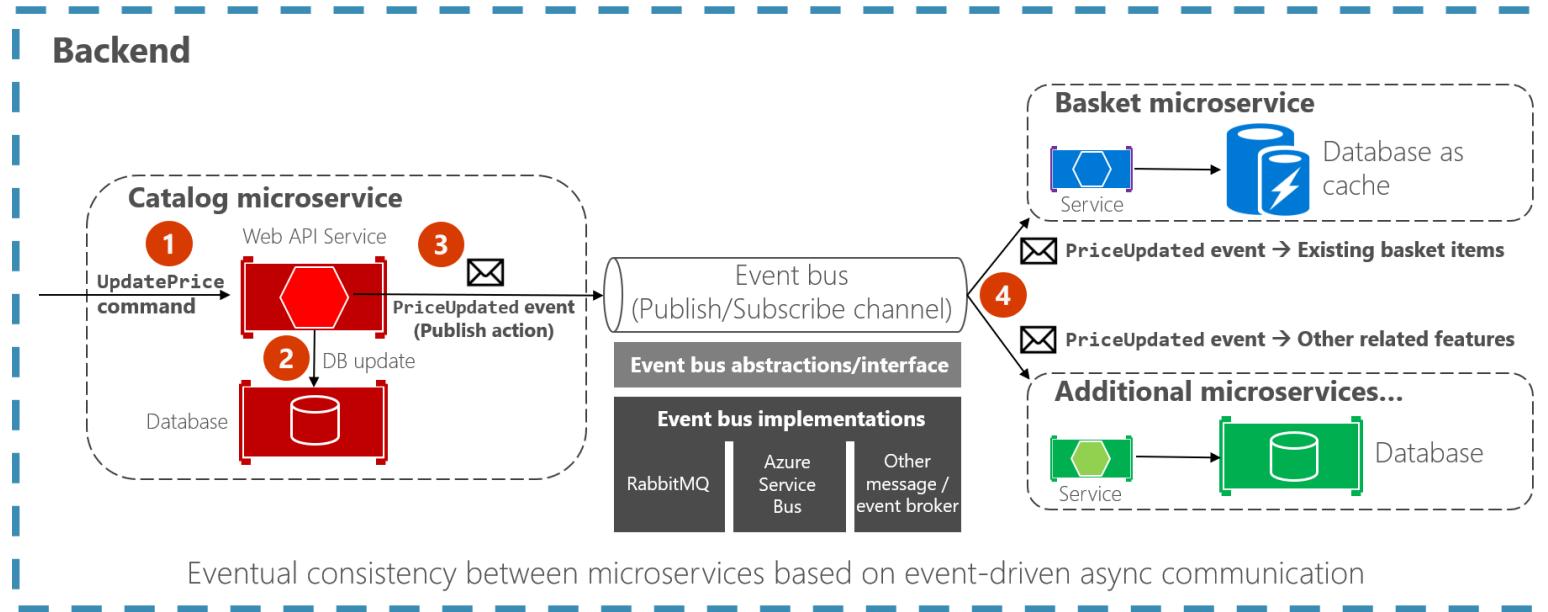
https://en.wikipedia.org/wiki/Publish%E2%80%93subscribe_pattern

For microservices:

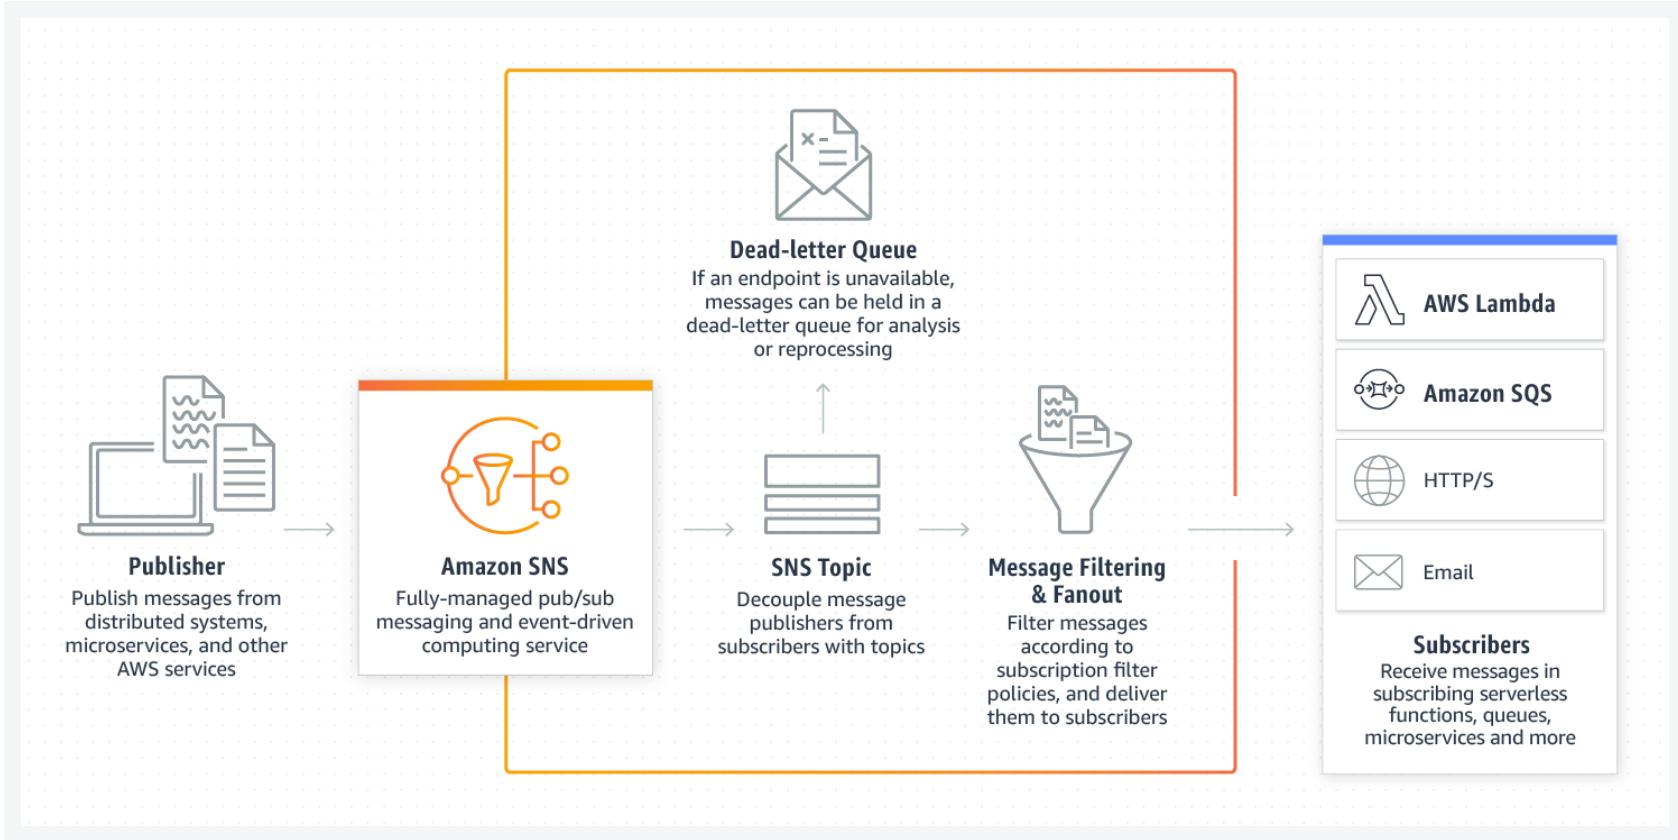
- It supports loose coupling for application evolution and agility.
- I can add and evolve services without modifying the base services.
- For example,
 - If creating a customer service **calls** "send an email."
 - I have to modify the original code to add support for SMS, webhook, etc.
 - With EDA, I just add or modify subscriptions.



Implementing asynchronous event-driven communication with an event bus

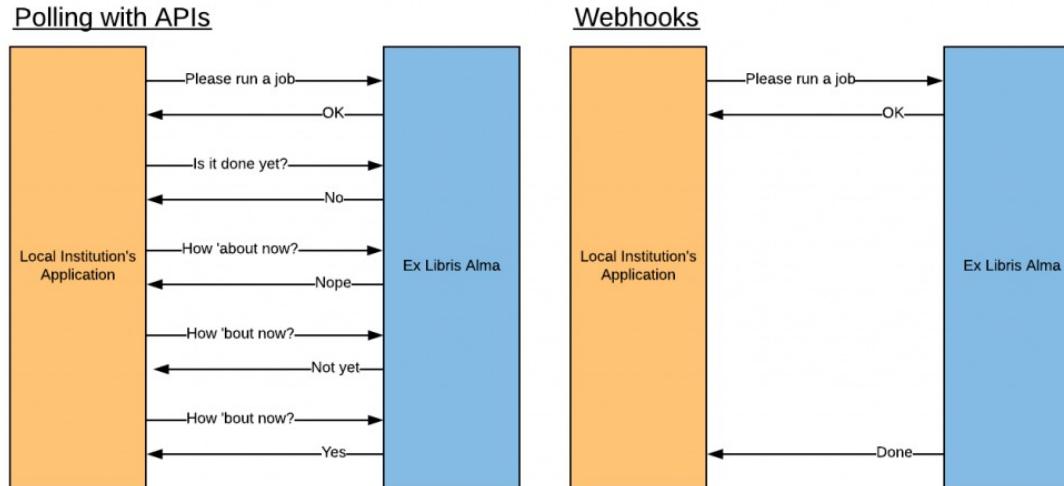


<https://docs.microsoft.com/en-us/dotnet/architecture/microservices/multi-container-microservice-net-applications/integration-event-based-microservice-communications>



Webhooks

- “A webhook in web development is a method of augmenting or altering the behavior of a web page or web application with custom callbacks. These callbacks may be maintained, modified, and managed by third-party users and developers who may not necessarily be affiliated with the originating website or application. The term “webhook” was coined by Jeff Lindsay in 2007 from the computer programming term hook.”



Function-as-a-Service

Serverless and Function-as-a-Service

- **Serverless computing** is a [cloud computing execution model](#) in which the cloud provider runs the [server](#), and dynamically manages the allocation of machine resources. Pricing is based on the actual amount of resources consumed by an application, rather than on pre-purchased units of capacity.^[1] It can be a form of [utility computing](#).

Serverless computing can simplify the process of [deploying code](#) into production. Scaling, capacity planning and maintenance operations may be hidden from the developer or operator. Serverless code can be used in conjunction with code deployed in traditional styles, such as [microservices](#). Alternatively, applications can be written to be purely serverless and use no provisioned servers at all.^[2] This should not be confused with computing or networking models that do not require an actual server to function, such as [peer-to-peer](#) (P2P)."
- **Function as a service (FaaS)** is a category of [cloud computing services](#) that provides a [platform](#) allowing customers to develop, run, and manage application functionalities without the complexity of building and maintaining the infrastructure typically associated with developing and launching an app.^[1] Building an application following this model is one way of achieving a "[serverless](#)" architecture, and is typically used when building [microservices](#) applications."
- That was baffling:
 - IaaS, CaaS – You control the SW stack and the cloud provides (virtual) HW.
 - PaaS – The cloud hides the lower layer SW and provides an application container (e.g. Flask) with “Your code goes here.” You are aware of container.
 - FaaS – The cloud provides the container, and you implement functions (corresponding to routes in REST).

Serverless and Function-as-a-Service

Private Cloud	IaaS	PaaS	FaaS	SaaS
	Infrastructure as a Service	Platform as a Service	Function as a Service	Software as a Service
Function	Function	Function	Function	Function
Application	Application	Application	Application	Application
Runtime	Runtime	Runtime	Runtime	Runtime
Operating System	Operating System	Operating System	Operating System	Operating System
Virtualization	Virtualization	Virtualization	Virtualization	Virtualization
Server	Server	Server	Server	Server
Storage	Storage	Storage	Storage	Storage
Networking	Networking	Networking	Networking	Networking

<https://medium.com/@tanmayct/serverless-architecture-function-as-a-service-19e127b8c990>

Managed by the customer



Managed by the provider



What is Serverless Good/Not Good For ... ?

Serverless is **good** for
*short-running
stateless
event-driven*



- Microservices
- Mobile Backends
- Bots, ML Inferencing
- IoT
- Modest Stream Processing
- Service integration

Serverless is **not good** for
*long-running
stateful
number crunching*

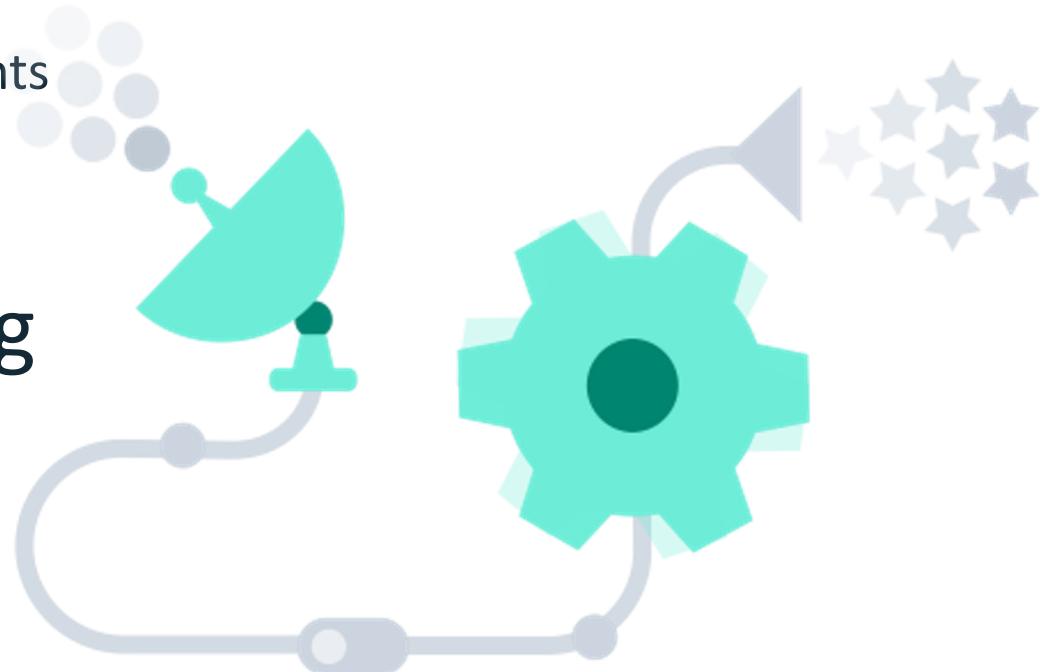


- Databases
- Deep Learning Training
- Heavy-Duty Stream Analytics
- Numerical Simulation
- Video Streaming

What triggers code execution?

Runs code **in response** to events

Event-programming
model



(Some) Current Platforms for Serverless



IBM Cloud
Functions



Azure
Functions



Red-Hat



Google
Functions



Kubernetes



Let's Build a Lambda Function

- We are going to do it manually to understand the process/steps.
- There are several more complete approaches:
 - Pipelines
 - Serverless Framework
(<https://www.serverless.com/framework/docs/providers/aws/guide/intro>)
 - AWS Toolkit for PyCharm
(<https://aws.amazon.com/pycharm/>)
- We will explore some of these

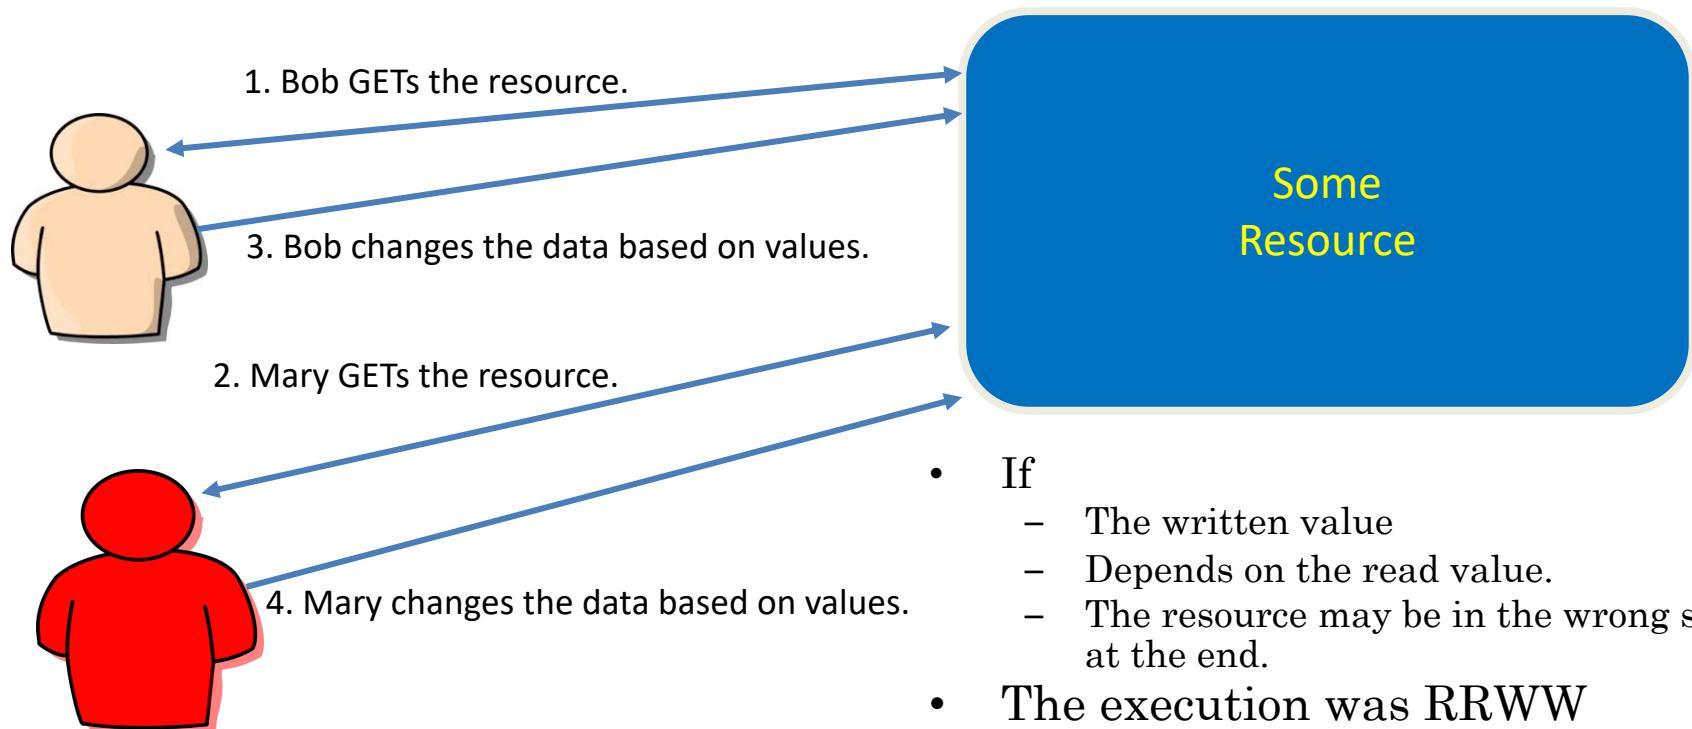
Putting Some Pieces Together

Projects

- /Users/donaldferguson/Dropbox/00NewProjects/SimpleMiddlerware
- /Users/donaldferguson/Dropbox/00NewProjects/slack-lambda/hello_world/slack_connector.py
- E6156-Student-Event Lambda function via Lambda Console
- /Users/donaldferguson/Dropbox/00NewProjects/google_login

ETag, Conditional Update

REST and Isolation: Read then Update Conflict



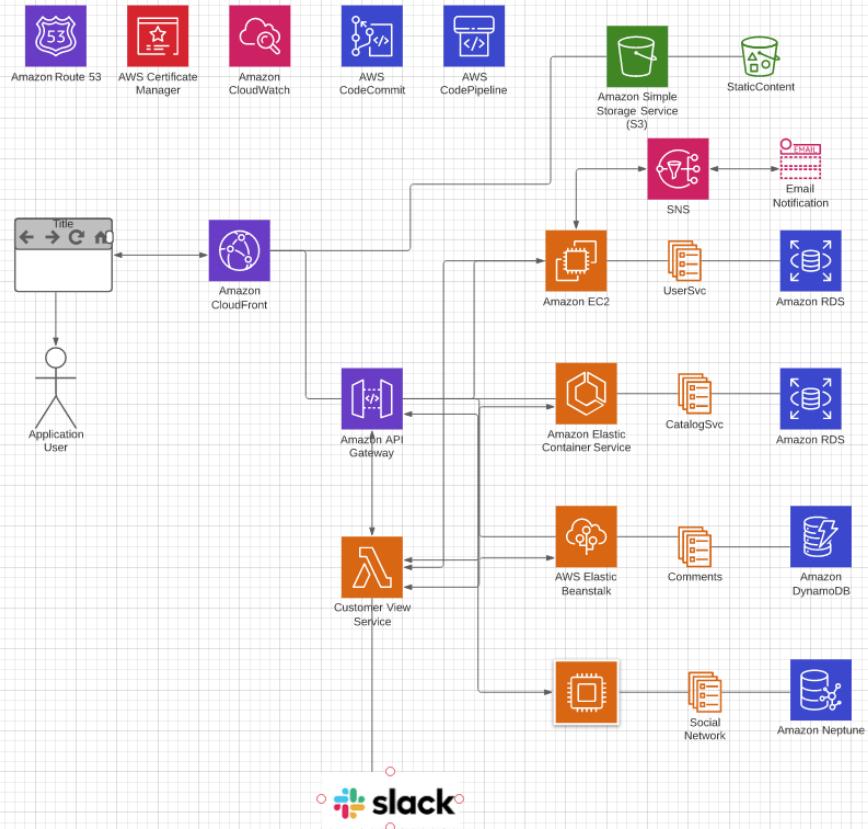
Isolation/Concurrency Control

This requires conversation state,
which would violate the REST
Stateless principle.

- There are two basic approaches to implementing isolation
 - Locking/Pessimistic, e.g. cursor isolation
 - Optimistic: Before committing, each transaction verifies that no other transaction has modified the data it has read.
- How does this work in REST?
 - The server maintains an ETag (Entity Tag) for each resource.
 - Every time a resource's state changes, the server computes a new ETag.
 - The server includes the ETag in the header when returning data to the client.
 - The server may *optionally* maintain and return "Last Modified" time for resources.
- Semantics on updates
 - If-Match – value of a previous calls ETag response used with a PUT or DELETE. Server should only act if nobody else has modified the resource since you last fetched it. Otherwise provides a 412.
 - If-Modified-Since – value of a previous Last-Modified response used with a GET. Server should only provide a response if the resource was modified since the timestamp submitted. Use in conjunction with If-None-Match in case the change occurred within the same second. Otherwise provide a 304.
 - If-None-Match – value of a previous calls ETag response used with a GET. Server should only provide a response if the ETag doesn't match, i.e. the resource has been altered. Otherwise provide a 304.
 - If-Unmodified-Since – value of a previous Last-Modified response used with a PUT or DELETE. Server should only act if nobody else has modified the resource since you last fetched it. Otherwise provides a 412

Hands-On CloudFront DNS, Certificate, HTTPS

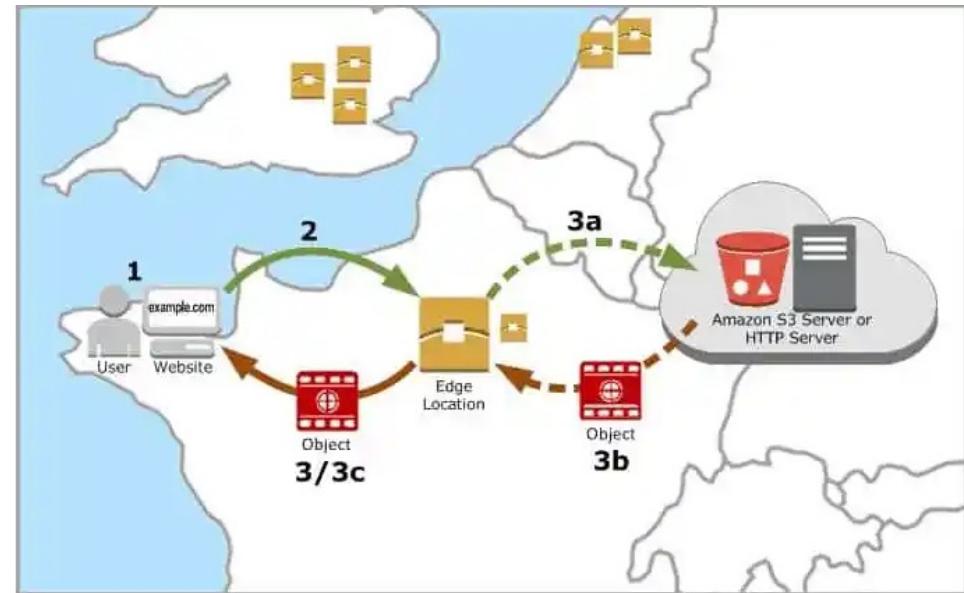
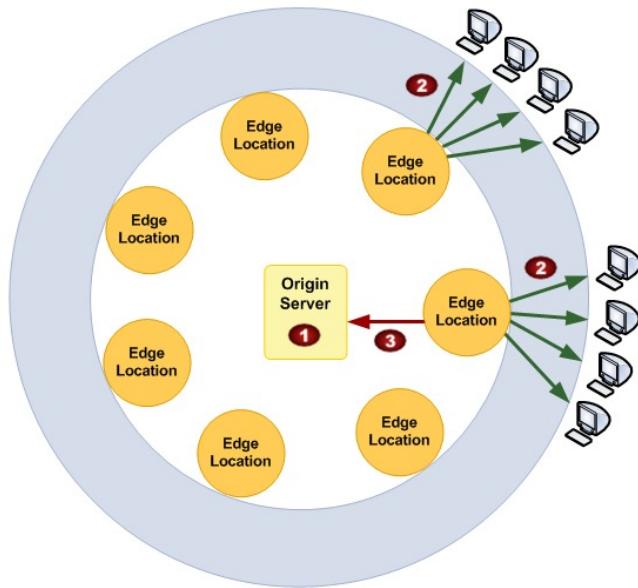
Overall Architecture



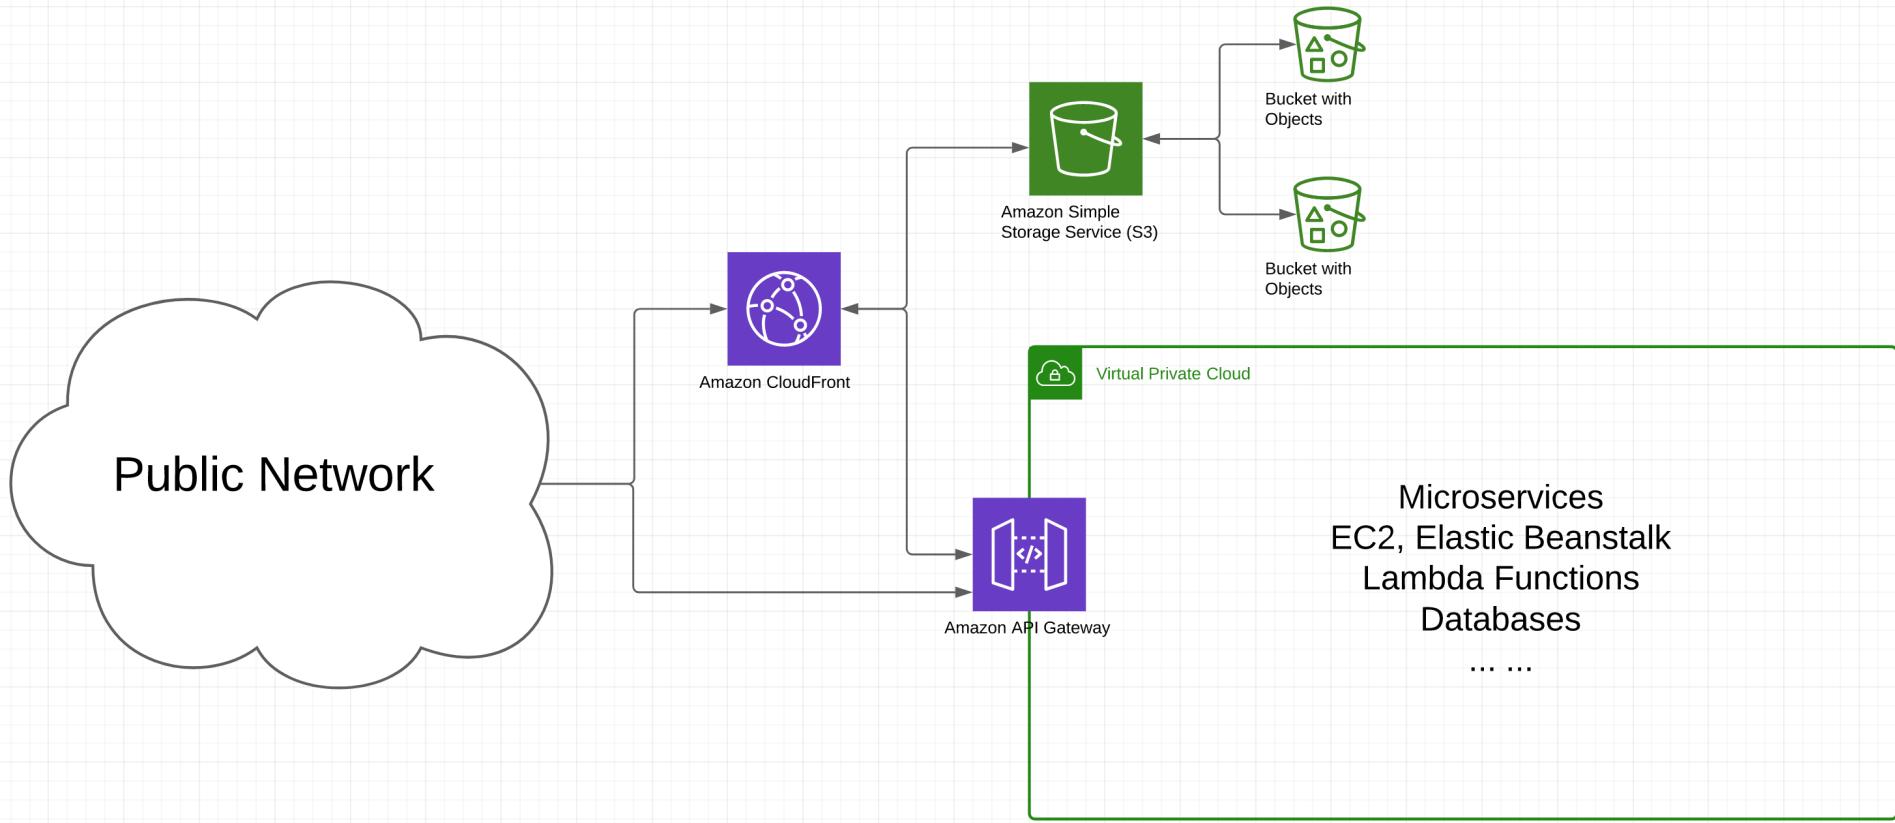
- There may be multiple
 - Microservices
 - Web UI projects/content
- CloudFront
 - Multiple S3 content → one site.
 - Forward API requests to API GW
- API GW manages/monitors APIs
- Networking/HTTPS requires
 - Domain, DNS
 - Certificates

CloudFront

“A content delivery network, or content distribution network (CDN), is a geographically distributed network of proxy servers and their data centers. The goal is to provide high availability and performance by distributing the service spatially relative to end users.”
https://en.wikipedia.org/wiki/Content_delivery_network



Big Picture

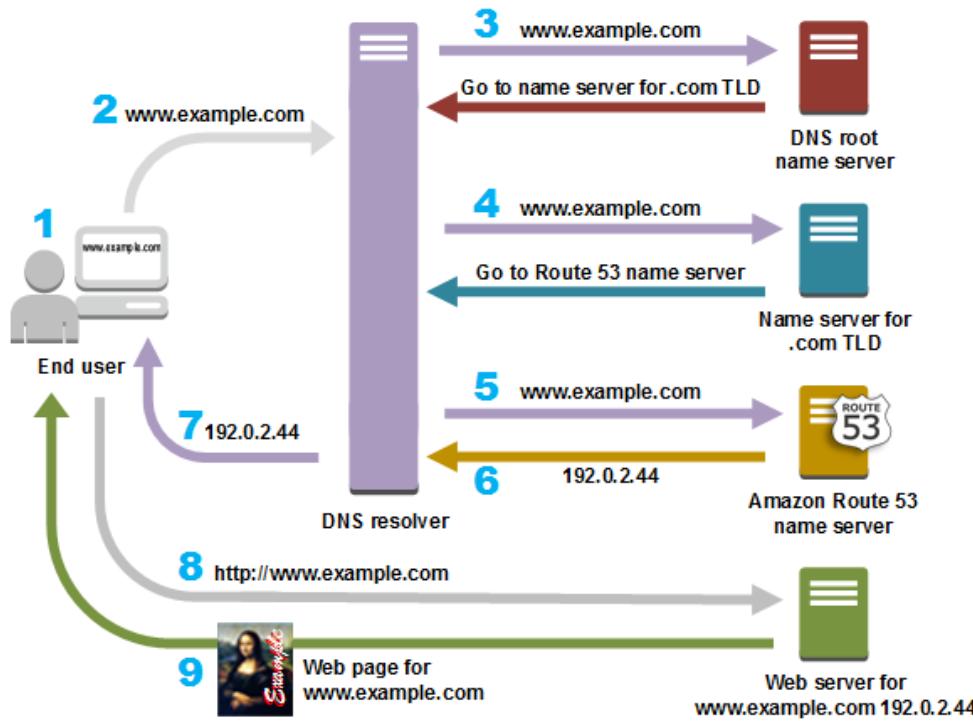


Certificates and HTTPS

How does HTTPS work: SSL explained

This presumes that SSL has already been issued by SSL issuing authority.





Things to Demo

- CloudFront:
 - Distributions, Origins, Behavior
 - Monitoring, caching, alarms,
- Certificates
 - AWS implementation
 - Use in CloudFront and elsewhere
 - Visibility in browser, HTTPs
- Route 53
 - Proving ownership
 - Records